Chapter 4bb

4.7 Summarizing data

An important part of exploratory data analysis is summarizing data. The average and standard deviation are two examples of widely used summary statistics. More informative summaries can often be achieved by first splitting data into groups. In this section, we cover two new dplyr verbs that make these computations easier: summarize and group_by. We learn to access resulting values using the pull function.

4.7.1 summarize

The summarize function in dplyr provides a way to compute summary statistics with intuitive and readable code. We start with a simple example based on heights. The heights dataset includes heights and sex reported by students in an in-class survey.

```
library(dplyr)
```

```
## Warning:   'dplyr' R   4.1.1
```

```
##
##          : 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##     filter, lag
```

```
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
```

```
library(dslabs)
```

```
## Warning:   'dslabs' R   4.1.1
```

```
data(heights)
```

The following code computes the average and standard deviation for females:

```
s <- heights %>%
  filter(sex == "Female") %>%
  summarize(average = mean(height), standard_deviation = sd(height))
s
```

```
##    average standard_deviation
## 1 64.93942         3.760656
```

This takes our original data table as input, filters it to keep only females, and then produces a new summarized table with just the average and the standard deviation of heights. We get to choose the names of the columns of the resulting table. For example, above we decided to use average and standard_deviation, but we could have used other names just the same.

Because the resulting table stored in s is a data frame, we can access the components with the accessor $:

```
s$average
```

```
## [1] 64.93942
```

```
s$standard_deviation
```

```
## [1] 3.760656
```

As with most other dplyr functions, summarize is aware of the variable names and we can use them directly. So when inside the call to the summarize function we write mean(height), the function is accessing the column with the name "height" and then computing the average of the resulting numeric vector. We can compute any other summary that operates on vectors and returns a single value.

For another example of how we can use the summarize function, let's compute the average murder rate for the United States. Remember our data table includes total murders and population size for each state and we have already used dplyr to add a murder rate column:

```
murders <- murders %>% mutate(rate = total/population*100000)
```

Remember that the US murder rate is not the average of the state murder rates:

```
summarize(murders, mean(rate))
```

```
##   mean(rate)
## 1   2.779125
```

This is because in the computation above the small states are given the same weight as the large ones. The US murder rate is the total number of murders in the US divided by the total US population. So the correct computation is:

```
us_murder_rate <- murders %>%
  summarize(rate = sum(total) / sum(population) * 100000)
us_murder_rate
```

```
##       rate
## 1 3.034555
```

This computation counts larger states proportionally to their size which results in a larger value.

4.7.2 Multiple summaries

Suppose we want three summaries from the same variable such as the median, minimum, and maximum heights. We can use summarize like this:

But we can obtain these three values with just one line using the quantile function: quantile(x, c(0.5, 0, 1)) returns the median (50th percentile), the min (0th percentile), and max (100th percentile) of the vector x. We can use it with summarize like this:

```
heights %>%
  filter(sex == "Female") %>%
  summarize(median_min_max = quantile(height, c(0.5, 0, 1)))
```

```
##    median_min_max
## 1        64.98031
## 2        51.00000
## 3        79.00000
```

However, notice that the summaries are returned in a row each. To obtain the results in different columns, we have to define a function that returns a data frame like this:

```
median_min_max <- function(x){
  qs <- quantile(x, c(0.5, 0, 1))
  data.frame(median = qs[1], minimum = qs[2], maximum = qs[3])
}
heights %>%
  filter(sex == "Female") %>%
  summarize(median_min_max(height))
```

```
##      median minimum maximum
## 1 64.98031      51      79
```

In the next section we learn how useful this approach can be when summarizing by group.

4.7.3 Group then summarize with group_by

A common operation in data exploration is to first split data into groups and then compute summaries for each group. For example, we may want to compute the average and standard deviation for men's and women's heights separately. The group_by function helps us do this.

If we type this:

```
heights %>% group_by(sex)
```

```
## # A tibble: 1,050 x 2
## # Groups:   sex [2]
##    sex    height
##    <fct>   <dbl>
## 1 Male       75
## 2 Male       70
## 3 Male       68
## 4 Male       74
```

```
##  5 Male        61
##  6 Female      65
##  7 Female      66
##  8 Female      62
##  9 Female      66
## 10 Male        67
## # ... with 1,040 more rows
```

The result does not look very different from heights, except we see Groups: sex [2] when we print the object. Although not immediately obvious from its appearance, this is now a special data frame called a grouped data frame, and dplyr functions, in particular summarize, will behave differently when acting on this object. Conceptually, you can think of this table as many tables, with the same columns but not necessarily the same number of rows, stacked together in one object. When we summarize the data after grouping, this is what happens:

```
heights %>%
  group_by(sex) %>%
  summarize(average = mean(height), standard_deviation = sd(height))
```

```
## # A tibble: 2 x 3
##   sex     average standard_deviation
##   <fct>     <dbl>              <dbl>
## 1 Female     64.9               3.76
## 2 Male       69.3               3.61
```

The summarize function applies the summarization to each group separately.

For another example, let's compute the median, minimum, and maximum murder rate in the four regions of the country using the median_min_max defined above:

```
murders %>%
  group_by(region) %>%
  summarize(median_min_max(rate))
```

```
## # A tibble: 4 x 4
##   region        median minimum maximum
##   <fct>          <dbl>   <dbl>   <dbl>
## 1 Northeast       1.80   0.320    3.60
## 2 South           3.40   1.46    16.5
## 3 North Central   1.97   0.595    5.36
## 4 West            1.29   0.515    3.63
```

4.8 pull

The us_murder_rate object defined above represents just one number. Yet we are storing it in a data frame:

```
class(us_murder_rate)
```

```
## [1] "data.frame"
```

since, as most dplyr functions, summarize always returns a data frame.

This might be problematic if we want to use this result with functions that require a numeric value. Here we show a useful trick for accessing values stored in data when using pipes: when a data object is piped that object and its columns can be accessed using the pull function. To understand what we mean take a look at this line of code:

```
us_murder_rate %>% pull(rate)
```

```
## [1] 3.034555
```

This returns the value in the rate column of us_murder_rate making it equivalent to us_murder_rate$rate.

To get a number from the original data table with one line of code we can type:

```
us_murder_rate <- murders %>%
  summarize(rate = sum(total) / sum(population) * 100000) %>%
  pull(rate)
us_murder_rate
```

```
## [1] 3.034555
```

which is now a numeric:

```
class(us_murder_rate)
```

```
## [1] "numeric"
```

4.9 Sorting data frames

When examining a dataset, it is often convenient to sort the table by the different columns. We know about the order and sort function, but for ordering entire tables, the dplyr function arrange is useful. For example, here we order the states by population size:

```
murders %>%
  arrange(population) %>%
  head()
```

```
##                  state abb        region population total        rate
## 1              Wyoming  WY          West     563626     5   0.8871131
## 2 District of Columbia  DC         South     601723    99  16.4527532
## 3              Vermont  VT     Northeast     625741     2   0.3196211
## 4         North Dakota  ND North Central     672591     4   0.5947151
## 5               Alaska  AK          West     710231    19   2.6751860
## 6         South Dakota  SD North Central     814180     8   0.9825837
```

With arrange we get to decide which column to sort by. To see the states by murder rate, from lowest to highest, we arrange by rate instead:

```
murders %>%
  arrange(rate) %>%
  head()
```

```
##            state abb       region population total      rate
## 1       Vermont  VT    Northeast     625741     2 0.3196211
## 2 New Hampshire  NH    Northeast    1316470     5 0.3798036
## 3        Hawaii  HI         West    1360301     7 0.5145920
## 4  North Dakota  ND North Central    672591     4 0.5947151
## 5          Iowa  IA North Central   3046355    21 0.6893484
## 6         Idaho  ID         West    1567582    12 0.7655102
```

Note that the default behavior is to order in ascending order. In dplyr, the function desc transforms a vector so that it is in descending order. To sort the table in descending order, we can type:

```
murders %>%
  arrange(desc(rate))
```

```
##                   state abb       region population total        rate
## 1  District of Columbia  DC        South     601723    99 16.4527532
## 2             Louisiana  LA        South    4533372   351  7.7425810
## 3              Missouri  MO North Central   5988927   321  5.3598917
## 4              Maryland  MD        South    5773552   293  5.0748655
## 5        South Carolina  SC        South    4625364   207  4.4753235
## 6              Delaware  DE        South     897934    38  4.2319369
## 7              Michigan  MI North Central   9883640   413  4.1786225
## 8           Mississippi  MS        South    2967297   120  4.0440846
## 9               Georgia  GA        South    9920000   376  3.7903226
## 10              Arizona  AZ         West    6392017   232  3.6295273
## 11         Pennsylvania  PA    Northeast   12702379   457  3.5977513
## 12            Tennessee  TN        South    6346105   219  3.4509357
## 13              Florida  FL        South   19687653   669  3.3980688
## 14           California  CA         West   37253956  1257  3.3741383
## 15           New Mexico  NM         West    2059179    67  3.2537239
## 16                Texas  TX        South   25145561   805  3.2013603
## 17             Arkansas  AR        South    2915918    93  3.1893901
## 18             Virginia  VA        South    8001024   250  3.1246001
## 19               Nevada  NV         West    2700551    84  3.1104763
## 20       North Carolina  NC        South    9535483   286  2.9993237
## 21             Oklahoma  OK        South    3751351   111  2.9589340
## 22             Illinois  IL North Central  12830632   364  2.8369608
## 23              Alabama  AL        South    4779736   135  2.8244238
## 24           New Jersey  NJ    Northeast    8791894   246  2.7980319
## 25          Connecticut  CT    Northeast    3574097    97  2.7139722
## 26                 Ohio  OH North Central  11536504   310  2.6871225
## 27               Alaska  AK         West     710231    19  2.6751860
## 28             Kentucky  KY        South    4339367   116  2.6732010
## 29             New York  NY    Northeast   19378102   517  2.6679599
## 30               Kansas  KS North Central   2853118    63  2.2081106
## 31              Indiana  IN North Central   6483802   142  2.1900730
## 32        Massachusetts  MA    Northeast    6547629   118  1.8021791
## 33             Nebraska  NE North Central   1826341    32  1.7521372
```

6

```
## 34         Wisconsin  WI North Central  5686986   97 1.7056487
## 35       Rhode Island  RI     Northeast  1052567   16 1.5200933
## 36      West Virginia  WV         South  1852994   27 1.4571013
## 37         Washington  WA          West  6724540   93 1.3829942
## 38           Colorado  CO          West  5029196   65 1.2924531
## 39            Montana  MT          West   989415   12 1.2128379
## 40          Minnesota  MN North Central  5303925   53 0.9992600
## 41       South Dakota  SD North Central   814180    8 0.9825837
## 42             Oregon  OR          West  3831074   36 0.9396843
## 43            Wyoming  WY          West   563626    5 0.8871131
## 44               Maine ME     Northeast  1328361   11 0.8280881
## 45               Utah  UT          West  2763885   22 0.7959810
## 46              Idaho  ID          West  1567582   12 0.7655102
## 47               Iowa  IA North Central  3046355   21 0.6893484
## 48       North Dakota  ND North Central   672591    4 0.5947151
## 49             Hawaii  HI          West  1360301    7 0.5145920
## 50      New Hampshire  NH     Northeast  1316470    5 0.3798036
## 51            Vermont  VT     Northeast   625741    2 0.3196211
```

4.9.1 Nested sorting

If we are ordering by a column with ties, we can use a second column to break the tie. Similarly, a third column can be used to break ties between first and second and so on. Here we order by region, then within region we order by murder rate:

```
murders %>%
  arrange(region, rate) %>%
  head()
```

```
##             state abb    region population total      rate
## 1       Vermont  VT Northeast     625741     2 0.3196211
## 2 New Hampshire  NH Northeast    1316470     5 0.3798036
## 3         Maine  ME Northeast    1328361    11 0.8280881
## 4  Rhode Island  RI Northeast    1052567    16 1.5200933
## 5 Massachusetts  MA Northeast    6547629   118 1.8021791
## 6      New York  NY Northeast   19378102   517 2.6679599
```

4.9.2 The top n

In the code above, we have used the function head to avoid having the page fill up with the entire dataset. If we want to see a larger proportion, we can use the top_n function. This function takes a data frame as it's first argument, the number of rows to show in the second, and the variable to filter by in the third. Here is an example of how to see the top 5 rows:

```
murders %>% top_n(5, rate)
```

```
##                 state abb       region population total      rate
## 1 District of Columbia  DC        South     601723    99 16.452753
## 2            Louisiana  LA        South    4533372   351  7.742581
## 3             Maryland  MD        South    5773552   293  5.074866
## 4             Missouri  MO North Central  5988927   321  5.359892
## 5       South Carolina  SC        South    4625364   207  4.475323
```

Note that rows are not sorted by rate, only filtered. If we want to sort, we need to use arrange. Note that if the third argument is left blank, top_n filters by the last column.

4.10 Exercises

For these exercises, we will be using the data from the survey collected by the United States National Center for Health Statistics (NCHS). This center has conducted a series of health and nutrition surveys since the 1960's. Starting in 1999, about 5,000 individuals of all ages have been interviewed every year and they complete the health examination component of the survey. Part of the data is made available via the NHANES package. Once you install the NHANES package, you can load the data like this:

```
library(NHANES)
```

```
## Warning:    'NHANES' R   4.1.1
```

```
data(NHANES)
```

The NHANES data has many missing values. The mean and sd functions in R will return NA if any of the entries of the input vector is an NA. Here is an example:

```
library(dslabs)
data(na_example)
mean(na_example)
```

```
## [1] NA
```

```
sd(na_example)
```

```
## [1] NA
```

To ignore the NAs we can use the na.rm argument:

```
mean(na_example, na.rm = TRUE)
```

```
## [1] 2.301754
```

```
sd(na_example, na.rm = TRUE)
```

```
## [1] 1.22338
```

Let's now explore the NHANES data.

1. We will provide some basic facts about blood pressure. First let's select a group to set the standard. We will use 20-to-29-year-old females. AgeDecade is a categorical variable with these ages. Note that the category is coded like " 20-29", with a space in front! What is the average and standard deviation of systolic blood pressure as saved in the BPSysAve variable? Save it to a variable called ref.

```
ref <- NHANES %>% filter(Gender=="female" & AgeDecade==" 20-29") %>% summarize(average = mean(BPSysAve,
ref
```

```
## # A tibble: 1 x 2
##   average standard_deviation
##     <dbl>              <dbl>
## 1    108.               10.1
```

2. Using a pipe, assign the average to a numeric variable ref_avg. Hint: Use the code similar to above and then pull.

```
ref_avg <- NHANES %>% filter(Gender=="female" & AgeDecade==" 20-29") %>% summarize(average = mean(BPSys
ref_avg
```

```
## [1] 108.4224
```

3. Now report the min and max values for the same group.

```
NHANES %>% filter(Gender=="female" & AgeDecade==" 20-29") %>% summarize(min(BPSysAve,na.rm=TRUE)) %>% pu
```

```
## [1] 84
```

```
NHANES %>% filter(Gender=="female" & AgeDecade==" 20-29") %>% summarize(max(BPSysAve,na.rm=TRUE)) %>% pu
```

```
## [1] 179
```

4. Compute the average and standard deviation for females, but for each age group separately rather than a selected decade as in question 1. Note that the age groups are defined by AgeDecade. Hint: rather than filtering by age and gender, filter by Gender and then use group_by.

```
NHANES %>% filter(Gender=="female") %>% group_by(AgeDecade) %>% summarize(average=mean(BPSysAve,na.rm=TP
```

```
## # A tibble: 9 x 3
##   AgeDecade average standard_deviation
##   <fct>       <dbl>              <dbl>
## 1 " 0-9"       100.               9.07
## 2 " 10-19"     104.               9.46
## 3 " 20-29"     108.              10.1
## 4 " 30-39"     111.              12.3
## 5 " 40-49"     115.              14.5
## 6 " 50-59"     122.              16.2
## 7 " 60-69"     127.              17.1
## 8 " 70+"       134.              19.8
## 9  <NA>        142.              22.9
```

5. Repeat exercise 4 for males.

```
NHANES %>% filter(Gender=="male") %>% group_by(AgeDecade) %>% summarize(average=mean(BPSysAve,na.rm=TRU
```

```
## # A tibble: 9 x 3
##    AgeDecade average standard_deviation
##    <fct>      <dbl>             <dbl>
## 1 " 0-9"      97.4              8.32
## 2 " 10-19"   110.              11.2
## 3 " 20-29"   118.              11.3
## 4 " 30-39"   119.              12.3
## 5 " 40-49"   121.              14.0
## 6 " 50-59"   126.              17.8
## 7 " 60-69"   127.              17.5
## 8 " 70+"     130.              18.7
## 9  <NA>      136.              23.5
```

6. We can actually combine both summaries for exercises 4 and 5 into one line of code. This is because group_by permits us to group by more than one variable. Obtain one big summary table using group_by(AgeDecade, Gender).

```
NHANES %>% group_by(AgeDecade,Gender) %>% summarize(average=mean(BPSysAve,na.rm=TRUE),standard_deviatio
```

```
## `summarise()` has grouped output by 'AgeDecade'. You can override using the `.groups` argument.
```

```
## # A tibble: 18 x 4
## # Groups:   AgeDecade [9]
##    AgeDecade Gender average standard_deviation
##    <fct>     <fct>   <dbl>            <dbl>
##  1 " 0-9"    female  100.              9.07
##  2 " 0-9"    male     97.4             8.32
##  3 " 10-19"  female  104.              9.46
##  4 " 10-19"  male    110.             11.2
##  5 " 20-29"  female  108.             10.1
##  6 " 20-29"  male    118.             11.3
##  7 " 30-39"  female  111.             12.3
##  8 " 30-39"  male    119.             12.3
##  9 " 40-49"  female  115.             14.5
## 10 " 40-49"  male    121.             14.0
## 11 " 50-59"  female  122.             16.2
## 12 " 50-59"  male    126.             17.8
## 13 " 60-69"  female  127.             17.1
## 14 " 60-69"  male    127.             17.5
## 15 " 70+"    female  134.             19.8
## 16 " 70+"    male    130.             18.7
## 17  <NA>     female  142.             22.9
## 18  <NA>     male    136.             23.5
```

7. For males between the ages of 40-49, compare systolic blood pressure across race as reported in the Race1 variable. Order the resulting table from lowest to highest average systolic blood pressure.

```
NHANES %>% filter(Gender == "male" & AgeDecade == " 40-49") %>% group_by(Race1) %>% summarize(average_s
```

```
## # A tibble: 5 x 2
##   Race1    average_systolic_blood_pressure
##   <fct>                              <dbl>
## 1 White                               120.
## 2 Other                               120.
## 3 Hispanic                            122.
## 4 Mexican                             122.
## 5 Black                               126.
```

4.11 Tibbles

Tidy data must be stored in data frames. We introduced the data frame in Section 2.4.1 and have been using the murders data frame throughout the book. In Section 4.7.3 we introduced the group_by function, which permits stratifying data before computing summary statistics. But where is the group information stored in the data frame?

```
murders %>% group_by(region)
```

```
## # A tibble: 51 x 6
## # Groups:   region [4]
##    state              abb   region     population total  rate
##    <chr>              <chr> <fct>           <dbl> <dbl> <dbl>
##  1 Alabama            AL    South         4779736   135  2.82
##  2 Alaska             AK    West           710231    19  2.68
##  3 Arizona            AZ    West          6392017   232  3.63
##  4 Arkansas           AR    South         2915918    93  3.19
##  5 California         CA    West         37253956  1257  3.37
##  6 Colorado           CO    West          5029196    65  1.29
##  7 Connecticut        CT    Northeast     3574097    97  2.71
##  8 Delaware           DE    South          897934    38  4.23
##  9 District of Columbia DC  South          601723    99 16.5
## 10 Florida            FL    South        19687653   669  3.40
## # ... with 41 more rows
```

Notice that there are no columns with this information. But, if you look closely at the output above, you see the line A tibble followd by dimensions. We can learn the class of the returned object using:

```
murders %>% group_by(region) %>% class()
```

```
## [1] "grouped_df" "tbl_df"     "tbl"        "data.frame"
```

The tbl, pronounced tibble, is a special kind of data frame. The functions group_by and summarize always return this type of data frame. The group_by function returns a special kind of tbl, the grouped_df. We will say more about these later. For consistency, the dplyr manipulation verbs (select, filter, mutate, and arrange) preserve the class of the input: if they receive a regular data frame they return a regular data frame, while if they receive a tibble they return a tibble. But tibbles are the preferred format in the tidyverse and as a result tidyverse functions that produce a data frame from scratch return a tibble. For example, in Chapter 5 we will see that tidyverse functions used to import data create tibbles.

Tibbles are very similar to data frames. In fact, you can think of them as a modern version of data frames. Nonetheless there are three important differences which we describe next.

4.11.1 Tibbles display better

The print method for tibbles is more readable than that of a data frame. To see this, compare the outputs of typing murders and the output of murders if we convert it to a tibble. We can do this using as_tibble(murders). If using RStudio, output for a tibble adjusts to your window size. To see this, change the width of your R console and notice how more/less columns are shown.

4.11.2 Subsets of tibbles are tibbles

If you subset the columns of a data frame, you may get back an object that is not a data frame, such as a vector or scalar. For example:

```
class(murders[,4])
```

```
## [1] "numeric"
```

is not a data frame. With tibbles this does not happen:

```
class(as_tibble(murders)[,4])
```

```
## [1] "tbl_df"     "tbl"         "data.frame"
```

This is useful in the tidyverse since functions require data frames as input.

With tibbles, if you want to access the vector that defines a column, and not get back a data frame, you need to use the accessor $:

```
class(as_tibble(murders)$population)
```

```
## [1] "numeric"
```

A related feature is that tibbles will give you a warning if you try to access a column that does not exist. If we accidentally write Population instead of population this:

```
murders$Population
```

```
## NULL
```

returns a NULL with no warning, which can make it harder to debug. In contrast, if we try this with a tibble we get an informative warning:

```
as_tibble(murders)$Population
```

```
## Warning: Unknown or uninitialised column: `Population`.
```

```
## NULL
```

4.11.3 Tibbles can have complex entries

While data frame columns need to be vectors of numbers, strings, or logical values, tibbles can have more complex objects, such as lists or functions. Also, we can create tibbles with functions:

```r
tibble(id = c(1, 2, 3), func = c(mean, median, sd))
```

```
## # A tibble: 3 x 2
##      id func
##   <dbl> <list>
## 1     1 <fn>
## 2     2 <fn>
## 3     3 <fn>
```

4.11.4 Tibbles can be grouped

The function group_by returns a special kind of tibble: a grouped tibble. This class stores information that lets you know which rows are in which groups. The tidyverse functions, in particular the summarize function, are aware of the group information.

4.11.5 Create a tibble using tibble instead of data.frame

It is sometimes useful for us to create our own data frames. To create a data frame in the tibble format, you can do this by using the tibble function.

```r
grades <- tibble(names = c("John", "Juan", "Jean", "Yao"),
                 exam_1 = c(95, 80, 90, 85),
                 exam_2 = c(90, 85, 85, 90))
```

Note that base R (without packages loaded) has a function with a very similar name, data.frame, that can be used to create a regular data frame rather than a tibble.

```r
grades <- data.frame(names = c("John", "Juan", "Jean", "Yao"),
                     exam_1 = c(95, 80, 90, 85),
                     exam_2 = c(90, 85, 85, 90))
```

To convert a regular data frame to a tibble, you can use the as_tibble function.

```r
as_tibble(grades) %>% class()
```

```
## [1] "tbl_df"     "tbl"        "data.frame"
```

4.12 The dot operator

One of the advantages of using the pipe %>% is that we do not have to keep naming new objects as we manipulate the data frame. As a quick reminder, if we want to compute the median murder rate for states in the southern states, instead of typing:

```r
tab_1 <- filter(murders, region == "South")
tab_2 <- mutate(tab_1, rate = total / population * 10^5)
rates <- tab_2$rate
median(rates)
```

```
## [1] 3.398069
```

We can avoid defining any new intermediate objects by instead typing:

```
filter(murders, region == "South") %>%
  mutate(rate = total / population * 10^5) %>%
  summarize(median = median(rate)) %>%
  pull(median)
```

```
## [1] 3.398069
```

We can do this because each of these functions takes a data frame as the first argument. But what if we want to access a component of the data frame. For example, what if the pull function was not available and we wanted to access tab_2$rate? What data frame name would we use? The answer is the dot operator.

For example to access the rate vector without the pull function we could use

```
rates <-   filter(murders, region == "South") %>%
  mutate(rate = total / population * 10^5) %>%
  .$rate
median(rates)
```

```
## [1] 3.398069
```

4.13 The purrr package

In Section 3.5 we learned about the sapply function, which permitted us to apply the same function to each element of a vector. We constructed a function and used sapply to compute the sum of the first n integers for several values of n like this:

```
compute_s_n <- function(n){
  x <- 1:n
  sum(x)
}
n <- 1:25
s_n <- sapply(n, compute_s_n)
```

This type of operation, applying the same function or procedure to elements of an object, is quite common in data analysis. The purrr package includes functions similar to sapply but that better interact with other tidyverse functions. The main advantage is that we can better control the output type of functions. In contrast, sapply can return several different object types; for example, we might expect a numeric result from a line of code, but sapply might convert our result to character under some circumstances. purrr functions will never do this: they will return objects of a specified type or return an error if this is not possible.

The first purrr function we will learn is map, which works very similar to sapply but always, without exception, returns a list:

```
library(purrr)
```

```
## Warning:   'purrr' R   4.1.1
```

```
s_n <- map(n, compute_s_n)
class(s_n)
```

```
## [1] "list"
```

If we want a numeric vector, we can instead use map_dbl which always returns a vector of numeric values.

```
s_n <- map_dbl(n, compute_s_n)
class(s_n)
```

```
## [1] "numeric"
```

This produces the same results as the sapply call shown above.

A particularly useful purrr function for interacting with the rest of the tidyverse is map_df, which always returns a tibble data frame. However, the function being called needs to return a vector or a list with names. For this reason, the following code would result in a Argument 1 must have names error:

```
# s_n <- map_df(n, compute_s_n)
```

We need to change the function to make this work:

```
compute_s_n <- function(n){
  x <- 1:n
  tibble(sum = sum(x))
}
s_n <- map_df(n, compute_s_n)
```

The purrr package provides much more functionality not covered here. For more details you can consult this online resource.

4.14 Tidyverse conditionals

A typical data analysis will often involve one or more conditional operations. In Section 3.1 we described the ifelse function, which we will use extensively in this book. In this section we present two dplyr functions that provide further functionality for performing conditional operations.

4.14.1 case_when

The case_when function is useful for vectorizing conditional statements. It is similar to ifelse but can output any number of values, as opposed to just TRUE or FALSE. Here is an example splitting numbers into negative, positive, and 0:

```
x <- c(-2, -1, 0, 1, 2)
case_when(x < 0 ~ "Negative",
          x > 0 ~ "Positive",
          TRUE  ~ "Zero")
```

```
## [1] "Negative" "Negative" "Zero"     "Positive" "Positive"
```

A common use for this function is to define categorical variables based on existing variables. For example, suppose we want to compare the murder rates in four groups of states: New England, West Coast, South, and other. For each state, we need to ask if it is in New England, if it is not we ask if it is in the West Coast, if not we ask if it is in the South, and if not we assign other. Here is how we use case_when to do this:

```
murders %>%
  mutate(group = case_when(
    abb %in% c("ME", "NH", "VT", "MA", "RI", "CT") ~ "New England",
    abb %in% c("WA", "OR", "CA") ~ "West Coast",
    region == "South" ~ "South",
    TRUE ~ "Other")) %>%
  group_by(group) %>%
  summarize(rate = sum(total) / sum(population) * 10^5)
```

```
## # A tibble: 4 x 2
##   group       rate
##   <chr>      <dbl>
## 1 New England  1.72
## 2 Other        2.71
## 3 South        3.63
## 4 West Coast   2.90
```

4.14.2 between

A common operation in data analysis is to determine if a value falls inside an interval. We can check this using conditionals. For example, to check if the elements of a vector x are between a and b we can type

```
# x >= a & x <= b
```

However, this can become cumbersome, especially within the tidyverse approach. The between function performs the same operation.

```
# between(x, a, b)
```

4.15 Exercises

1. Load the murders dataset. Which of the following is true?

```
data(murders)
murders
```

```
##                    state abb      region population total
## 1                Alabama  AL       South    4779736   135
## 2                 Alaska  AK        West     710231    19
## 3                Arizona  AZ        West    6392017   232
## 4               Arkansas  AR       South    2915918    93
## 5             California  CA        West   37253956  1257
## 6               Colorado  CO        West    5029196    65
## 7            Connecticut  CT   Northeast    3574097    97
## 8               Delaware  DE       South     897934    38
## 9   District of Columbia  DC       South     601723    99
```

16

```
## 10           Florida  FL         South  19687653  669
## 11           Georgia  GA         South   9920000  376
## 12            Hawaii  HI          West   1360301    7
## 13             Idaho  ID          West   1567582   12
## 14          Illinois  IL North Central  12830632  364
## 15           Indiana  IN North Central   6483802  142
## 16              Iowa  IA North Central   3046355   21
## 17            Kansas  KS North Central   2853118   63
## 18          Kentucky  KY         South   4339367  116
## 19         Louisiana  LA         South   4533372  351
## 20             Maine  ME     Northeast   1328361   11
## 21          Maryland  MD         South   5773552  293
## 22     Massachusetts  MA     Northeast   6547629  118
## 23          Michigan  MI North Central   9883640  413
## 24         Minnesota  MN North Central   5303925   53
## 25       Mississippi  MS         South   2967297  120
## 26          Missouri  MO North Central   5988927  321
## 27           Montana  MT          West    989415   12
## 28          Nebraska  NE North Central   1826341   32
## 29            Nevada  NV          West   2700551   84
## 30     New Hampshire  NH     Northeast   1316470    5
## 31        New Jersey  NJ     Northeast   8791894  246
## 32        New Mexico  NM          West   2059179   67
## 33          New York  NY     Northeast  19378102  517
## 34    North Carolina  NC         South   9535483  286
## 35      North Dakota  ND North Central    672591    4
## 36              Ohio  OH North Central  11536504  310
## 37          Oklahoma  OK         South   3751351  111
## 38            Oregon  OR          West   3831074   36
## 39      Pennsylvania  PA     Northeast  12702379  457
## 40      Rhode Island  RI     Northeast   1052567   16
## 41    South Carolina  SC         South   4625364  207
## 42      South Dakota  SD North Central    814180    8
## 43         Tennessee  TN         South   6346105  219
## 44             Texas  TX         South  25145561  805
## 45              Utah  UT          West   2763885   22
## 46           Vermont  VT     Northeast    625741    2
## 47          Virginia  VA         South   8001024  250
## 48        Washington  WA          West   6724540   93
## 49     West Virginia  WV         South   1852994   27
## 50         Wisconsin  WI North Central   5686986   97
## 51           Wyoming  WY          West    563626    5
```

  a. murders is in tidy format and is stored in a tibble.
  b. murders is in tidy format and is stored in a data frame.
  c. murders is not in tidy format and is stored in a tibble.
  d. murders is not in tidy format and is stored in a data frame.

answer : b

  2. Use as_tibble to convert the murders data table into a tibble and save it in an object called mur-
     ders_tibble.

17

```
murders_tibble <- as_tibble(murders)
murders_tibble
```

```
## # A tibble: 51 x 5
##    state              abb   region    population total
##    <chr>              <chr> <fct>          <dbl> <dbl>
##  1 Alabama            AL    South        4779736   135
##  2 Alaska             AK    West          710231    19
##  3 Arizona            AZ    West         6392017   232
##  4 Arkansas           AR    South        2915918    93
##  5 California         CA    West        37253956  1257
##  6 Colorado           CO    West         5029196    65
##  7 Connecticut        CT    Northeast    3574097    97
##  8 Delaware           DE    South         897934    38
##  9 District of Columbia DC  South         601723    99
## 10 Florida            FL    South       19687653   669
## # ... with 41 more rows
```

3. Use the group_by function to convert murders into a tibble that is grouped by region.

```
murders %>% group_by(region)
```

```
## # A tibble: 51 x 5
## # Groups:   region [4]
##    state              abb   region    population total
##    <chr>              <chr> <fct>          <dbl> <dbl>
##  1 Alabama            AL    South        4779736   135
##  2 Alaska             AK    West          710231    19
##  3 Arizona            AZ    West         6392017   232
##  4 Arkansas           AR    South        2915918    93
##  5 California         CA    West        37253956  1257
##  6 Colorado           CO    West         5029196    65
##  7 Connecticut        CT    Northeast    3574097    97
##  8 Delaware           DE    South         897934    38
##  9 District of Columbia DC  South         601723    99
## 10 Florida            FL    South       19687653   669
## # ... with 41 more rows
```

4. Write tidyverse code that is equivalent to this code:

```
exp(mean(log(murders$population)))
```

```
## [1] 3675209
```

Write it using the pipe so that each function is called without arguments. Use the dot operator to access the population. Hint: The code should start with murders %>%.

```
murders %>% .$population %>% log(.) %>% mean(.) %>% exp(.)
```

```
## [1] 3675209
```

5. Use the map_df to create a data frame with three columns named n, s_n, and s_n_2. The first column should contain the numbers 1 through 100. The second and third columns should each contain the sum of 1 through n with n the row number.

```r
n <- 1:100
compute_s_n <- function(n){
  x <- 1:n
  tibble(s_n=sum(x))
}
compute_s_n_2 <- function(n){
  x <- 1:n
  tibble(s_n_2=sum(x))
}
df5 <- map_df(n,compute_s_n) %>% mutate(n=n,map_df(n,compute_s_n_2))
df5[,c(2,1,3)]
```

```
## # A tibble: 100 x 3
##        n   s_n s_n_2
##    <int> <int> <int>
## 1      1     1     1
## 2      2     3     3
## 3      3     6     6
## 4      4    10    10
## 5      5    15    15
## 6      6    21    21
## 7      7    28    28
## 8      8    36    36
## 9      9    45    45
## 10    10    55    55
## # ... with 90 more rows
```