

# GENCODE Tutorial

## [Tutorial] Human Genome Annotation

Tutorial for Tidyverse (Chapter 4)

### 1. Introduction

#### 1.1. What is gene annotation?

Over the past years, we have learnt that there are a number of chromosomes and genes in our genome. Counting the number of chromosomes is fairly easy but students might find difficult to say how many genes we have in our genome. If you can get an answer for this, could you tell how many genes encode protein and how many do not?

To answer this question, we need to access the database for gene annotation. Gene annotation is the process of making nucleotide sequence meaningful - where genes are located? whether it is protein-coding or noncoding. If you would like to get an overview of gene annotation, please find this link.

One of well-known collaborative efforts in gene annotation is the GENCODE consortium. It is a part of the Encyclopedia of DNA Elements (The ENCODE project consortium) and aims to identify all gene features in the human genome using a combination of computational analysis, manual annotation, and experimental validation (Harrow et al. 2012). You might find another database for gene annotation, like RefSeq, CCDS, and need to understand differences between them.

Figure 1. Comparison of GENCODE and RefSeq gene annotation and the impact of reference geneset on variant effect prediction (Frankish et al. 2015). A) Mean number of alternatively spliced transcripts per multi-exon protein-coding locus B) Mean number of unique CDS per multi-exon protein-coding locus C) Mean number of unique (non-redundant) exons per multi-exon protein-coding locus D) Percentage genomic coverage of unique (non-redundant) exons at multi-exon protein-coding loci.

In this tutorial, we will access to gene annotation from the GENCODE consortium and explore genes and functional elements in our genome.

#### 1.2. Aims

What we will do with this dataset:

- Be familiar with gene annotation modality.
- Tidy data and create a table for your analysis.
- Apply tidyverse functions for data munging.

Please note that there is better solution for getting gene annotation in R if you use a biomart. Our tutorial is only designed to have a practice on tidyverse exercise.

## 2. Explore your data

### 2.1. Unboxing your dataset

This tutorial will use a gene annotation file from the GENCODE. You will need to download the file from the GENCODE. If you are using terminal, please download file using wget:

```
# Run from your terminal, not R console  
# wget ftp://ftp.ebi.ac.uk/pub/databases/genocode/Gencode_human/release_31/genocode.v31.basic.annotation.gtf.gz  
# Once you downloaded the file, you won't need to download it again. So please comment out the command
```

Once you download the file, you can print out the first few lines using the following bash command (we will learn UNIX commands later):

```
# Run from your terminal, not R console  
# gzcat genocode.v31.basic.annotation.gtf.gz | head -7
```

The file is the GTF file format, which you will find most commonly in gene annotation. Please read the file format thoroughly in the link above.

For the tutorial, we need to load two packages. If the package is not installed in your system, please install it.

- tidyverse, a package you have learnt from the chapter 5.
- readr, a package provides a fast and friendly way to read. Since the file genocode.v31.basic.annotation.gtf.gz is pretty large, you will need some function to load data quickly into your workspace. readr in a part of tidyverse, so you can just load tidyverse to use readr functions.

Let's load the GTF file into your workspace. We will use read\_delim function from the readr package. This is much faster loading than read.delim or read.csv from R base. However, please keep in mind that some parameters and output class for read\_delim are slightly different from them.

```
library(tidyverse)
```

```
## Warning:   'tidyverse' R    4.1.1
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr   0.3.4  
## v tibble  3.1.3      v dplyr   1.0.7  
## v tidyr   1.1.4      v stringr 1.4.0  
## v readr   2.0.2      v forcats 0.5.1
```

```
## Warning:   'ggplot2' R    4.1.1
```

```
## Warning:   'tidyr' R     4.1.1
```

```
## Warning:   'readr' R     4.1.1
```

```
## Warning:   'purrr' R     4.1.1
```

```
## Warning:      'dplyr' R      4.1.1

## Warning:      'forcats' R      4.1.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()      masks stats::lag()

d = read_delim('encode.v31.basic.annotation.gtf.gz',
               delim='\t', skip = 5, progress = F,
               col_names = F)

## Rows: 1756502 Columns: 9

## -- Column specification -----
## Delimiter: "\t"
## chr (7): X1, X2, X3, X6, X7, X8, X9
## dbl (2): X4, X5

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Can you find out what the parameters mean? Few things to note are:

- The GTF file contains the first few lines for comments (#). In general, the file contains description, provider, date, format.
- The GTF file does not have column names so you will need to assign 'FALSE' for col\_names.

This is sort of canonical way to load your dataset into R. However, we are using a GTF format, which is specific to gene annotation so we can use a package to specifically handle a GTF file.

Here I introduce the package rtracklayer. Let's install the package first.

```
if (!requireNamespace("BiocManager", quietly = TRUE))
  install.packages("BiocManager")

BiocManager::install("rtracklayer")

## Bioconductor version 3.13 (BiocManager 1.30.16), R 4.1.0 (2021-05-18)

## Warning: package(s) not installed when version(s) same as current; use `force = TRUE` to
## re-install: 'rtracklayer'

## Old packages: 'lubridate', 'stringi', 'tibble'
```

Then, now you can read the GTF file using this package. Then, you can check the class of the object d.

```
d = rtracklayer::import('gencode.v31.basic.annotation.gtf.gz')
class(d)
```

```
## [1] "GRanges"
## attr(,"package")
## [1] "GenomicRanges"
```

You will find out that this is GRanges class. This is from the package Genomic Range, specifically dealing with genomic datasets but we are not heading into this in this tutorial. So please find this information if you are serious on this.

We are converting d into a data frame as following:

```
d = d %>% as.data.frame()
```

Let's overview few lines from the data frame, and explore what you get in this object.

```
head(d)
```

```
##      seqnames start   end width strand source      type score phase
## 1      chr1 11869 14409  2541      + HAVANA      gene    NA     NA
## 2      chr1 11869 14409  2541      + HAVANA transcript NA     NA
## 3      chr1 11869 12227   359      + HAVANA      exon    NA     NA
## 4      chr1 12613 12721   109      + HAVANA      exon    NA     NA
## 5      chr1 13221 14409  1189      + HAVANA      exon    NA     NA
## 6      chr1 12010 13670  1661      + HAVANA transcript NA     NA
##      gene_id
## 1 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1 2
## 2 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1 2
## 3 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1 2
## 4 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1 2
## 5 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1 2
## 6 ENSG00000223972.5 transcribed_unprocessed_pseudogene DDX11L1 2
##      hgnc_id      havana_gene      transcript_id
## 1 HGNC:37102 OTTHUMG00000000961.2      <NA>
## 2 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 3 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 4 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 5 HGNC:37102 OTTHUMG00000000961.2 ENST00000456328.2
## 6 HGNC:37102 OTTHUMG00000000961.2 ENST00000450305.2
##      transcript_type transcript_name transcript_support_level
## 1      <NA>      <NA>      <NA>
## 2      lncRNA      DDX11L1-202      1
## 3      lncRNA      DDX11L1-202      1
## 4      lncRNA      DDX11L1-202      1
## 5      lncRNA      DDX11L1-202      1
## 6 transcribed_unprocessed_pseudogene DDX11L1-201      NA
##      tag      havana_transcript exon_number      exon_id      ont
## 1 <NA>      <NA>      <NA>      <NA>      <NA>
## 2 basic OTTHUMT00000362751.1      <NA>      <NA>      <NA>
## 3 basic OTTHUMT00000362751.1      1 ENSE000002234944.1      <NA>
## 4 basic OTTHUMT00000362751.1      2 ENSE000003582793.1      <NA>
## 5 basic OTTHUMT00000362751.1      3 ENSE000002312635.1      <NA>
```

```
## 6 basic OTTHUMT00000002844.2      <NA>      <NA> PGO:0000019
##   protein_id ccdsid
## 1      <NA>   <NA>
## 2      <NA>   <NA>
## 3      <NA>   <NA>
## 4      <NA>   <NA>
## 5      <NA>   <NA>
## 6      <NA>   <NA>
```

One thing you can find is that there is no columns in the data frame. Let's match which information is provided in columns. You can find the instruction page in the website ([link](#)).

Based on this, you can assign a name for 9 columns. One thing to remember is you should not use space for the column name. Spacing in the column name is actually working but not a good habit for your code. So please replace a space with underscore in the column name.

```
# Assign column names according to the GENCODE instruction.
cols = c('chrom', 'source', 'feature_type', 'start', 'end', 'score', 'strand', 'phase', 'info')
```

Now you can set up the column names into the `col_names` parameter, and load the file into a data frame.

```
d = read_delim('gencode.v31.basic.annotation.gtf.gz',
               delim='\t', skip = 5,
               progress = F,
               col_names = cols)
```

```
## Rows: 1756502 Columns: 9
```

```
## -- Column specification -----
## Delimiter: "\t"
## chr (7): chrom, source, feature_type, score, strand, phase, info
## dbl (2): start, end

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

You can find the column names are now all set.

```
head(d)
```

```
## # A tibble: 6 x 9
##   chrom source feature_type start   end score strand phase info
##   <chr> <chr>   <chr>      <dbl> <dbl> <chr> <chr>   <chr> <chr>
## 1 chr1  HAVANA  gene        11869 14409 .    +    .    "gene_id \"ENSG000000~
## 2 chr1  HAVANA  transcript   11869 14409 .    +    .    "gene_id \"ENSG000000~
## 3 chr1  HAVANA  exon        11869 12227 .    +    .    "gene_id \"ENSG000000~
## 4 chr1  HAVANA  exon        12613 12721 .    +    .    "gene_id \"ENSG000000~
## 5 chr1  HAVANA  exon        13221 14409 .    +    .    "gene_id \"ENSG000000~
## 6 chr1  HAVANA  transcript   12010 13670 .    +    .    "gene_id \"ENSG000000~
```

When you loaded the file, you see the message about the data class. You might want to overview this data.

```
summary(d)
```

```
##      chrom      source      feature_type      start
## Length:1756502 Length:1756502 Length:1756502 Min. :      577
## Class :character Class :character Class :character 1st Qu.: 32101517
## Mode :character Mode :character Mode :character Median : 61732754
##                                     Mean : 75288563
##                                     3rd Qu.:111760181
##                                     Max. :248936581
##      end      score      strand      phase
## Min. :      647 Length:1756502 Length:1756502 Length:1756502
## 1st Qu.: 32107331 Class :character Class :character Class :character
## Median : 61738373 Mode :character Mode :character Mode :character
## Mean : 75292632
## 3rd Qu.:111763007
## Max. :248937043
##      info
## Length:1756502
## Class :character
## Mode :character
##
##
##
```

## 2.2. How many feature types in the GENCODE dataset?

As instructed in the GENCODE website, the GENCODE dataset provides a range of annotations for the feature type. You can check feature types using `group_by`, and `count` or `table` function.

```
d %>% group_by(feature_type) %>% count(feature_type)
```

```
## # A tibble: 8 x 2
## # Groups:   feature_type [8]
##   feature_type      n
##   <chr>          <int>
## 1 CDS            567862
## 2 exon           744835
## 3 gene           60603
## 4 Selenocysteine    96
## 5 start_codon      57886
## 6 stop_codon       57775
## 7 transcript       108243
## 8 UTR             159202
```

```
# table(d$feature_type)
# 8 feature types from the dataset
```

How many feature types provided in the GENCODE? And how many items stored for each feature type? Please write down the number of feature types from the dataset. Also, if you are not familiar with these types, it would be good to put one or two sentences that can describe each type.

### 2.3. How many genes we have?

Let's count the number of genes in our genome. Since we know that the column `feature_type` contains rows with gene, which contains obviously annotations for genes. We might want to subset those rows from the data frame.

```
d1 = filter(d, feature_type == 'gene')
# d1 = d[d$feature_type == 'gene', ]
```

### 2.4. Ensembl, Havana and CCDS.

Gene annotation for the human genome is provided by multiple organizations with different gene annotation methods and strategy. This means that information can be varying by resources, and users need to understand heterogeneity inherent in annotation databases.

The GENCODE project utilizes two sources of gene annotation.

1. Havana: Manual gene annotation (detailed strategy in here)
2. Ensembl: Automatic gene annotation (detailed strategy in here)

It provides the combination of Ensembl/HAVANA gene set as the default gene annotation for the human genome. In addition, they also guarantee that all transcripts from the Consensus Coding Sequence (CCDS) set are present in the GENCODE gene set. The CCDS project is a collaborative effort to identify a core set of protein coding regions that are consistently annotated and of high quality. Initial results from the Consensus CDS (CCDS) project are now available through the appropriate Ensembl gene pages and from the CCDS project page at NCBI. The CCDS set is built by consensus among Ensembl, the National Center for Biotechnology Information (NCBI), and the HUGO Gene Nomenclature Committee (HGNC) for human (link).

Figure 2. Comparison of CCDS and Gencode (Source).

Right. Then now we count how many genes annotated with HAVANA and ENSEMBL.

```
d %>% group_by(source) %>% count(source)
```

```
## # A tibble: 2 x 2
## # Groups:   source [2]
##   source      n
##   <chr>    <int>
## 1 ENSEMBL 245185
## 2 HAVANA  1511317
```

### 2.5. do.call

Since the last column `info` contains a long string for multiple annotations, we will need to split it to extract each annotation. For example, the first line for transcript annotation looks like this:

```
# chr1    HAVANA    transcript    11869    14409    .    +    .    gene_id "ENSG00000223972.5"; transcr
```

If you would like to split `transcript_support_level` and create a new column, you can use `strsplit` function.

```
a = 'chr1    HAVANA    transcript    11869    14409    .    +    .    gene_id "ENSG00000223972.5"; trans
strsplit(a, 'transcript_support_level\\s+')

```

```
## [[1]]
## [1] "chr1    HAVANA    transcript    11869    14409    .    +    .    gene_id \"ENSG00000223972.5\"";
## [2] "1\"; hgnc_id \"HGNC:37102\"; tag \"basic\"; havana_gene \"OTTHUMG00000000961.2\"; havana_transc

```

After split the string, you can select the second item in the list ([[1]][2]).

```
strsplit(a, 'transcript_support_level\\s+')[[1]][2]

```

```
## [1] "1\"; hgnc_id \"HGNC:37102\"; tag \"basic\"; havana_gene \"OTTHUMG00000000961.2\"; havana_transc

```

You can find the 1 in the first position, which you will need to split again.

```
b = strsplit(a, 'transcript_support_level\\s+')[[1]][2]
strsplit(b, '\\')

```

```
## [[1]]
## [1] "1"                "; hgnc_id "        "HGNC:37102"
## [4] "; tag "           "basic"           "; havana_gene "
## [7] "OTTHUMG00000000961.2" "; havana_transcript " "OTTHUMT00000362751.1"
## [10] "; "

```

From this, you will get the first item in the list ([[1]][1]).

Now you would like to apply strsplit function across vectors. For this, do.call function can be easily implemented to strsplit over the vectors from one column. Let's try this.

```
head(do.call(rbind.data.frame, strsplit(a, 'transcript_support_level\\s+'))[[2]])

```

```
## [1] "1\"; hgnc_id \"HGNC:37102\"; tag \"basic\"; havana_gene \"OTTHUMG00000000961.2\"; havana_transc

```

Now you can write two lines of codes to process two steps we discussed above.

```
# First filter transcripts and create a data frame.
d2 <- d %>% filter(feature_type == 'transcript')

# Now apply the functions.
d2$transcript_support_level <- as.character(do.call(rbind.data.frame,
strsplit(d2$info, 'transcript_support_level\\s+'))[[2]])

d2$transcript_support_level <- as.character(do.call(rbind.data.frame,
strsplit(d2$transcript_support_level, '\\'))[[1]])

```

Now you can check the strsplit works.



```
head(d2$transcript_support_level)
```

```
## [1] "1" "NA" "NA" "NA" "5" "5"
```

You can use the same method to extract other annotations, like `gene_id`, `gene_name` etc.

### 3. Exercises

Here I list the questions for your activity. Please note that it is an exercise for tidyverse functions, which you will need to use in your code. In addition, you will need to write an one-line code for each question using pipe `%>%`.

For questions, you should read some information thoroughly, including:

- Gene biotype.
- 0 or 1 based annotation in GTF, BED format
- Why some features have 1 bp length?
- What is the meaning of zero-length exons in GENCODE? Also fun to have a review for microexons
- Transcript support level (TSL)

```
if (!requireNamespace("BiocManager", quietly = TRUE))
  install.packages("BiocManager")
BiocManager::install("rtracklayer")
```

```
## Bioconductor version 3.13 (BiocManager 1.30.16), R 4.1.0 (2021-05-18)
```

```
## Warning: package(s) not installed when version(s) same as current; use `force = TRUE` to
## re-install: 'rtracklayer'
```

```
## Old packages: 'lubridate', 'stringi', 'tibble'
```

```
d = rtracklayer::import('gencode.v31.basic.annotation.gtf.gz')
d = d %>% as.data.frame()
```

#### 3.1. Annotation of transcripts in our genome

1. Computes the number of transcripts per gene. What is the mean number of transcripts per gene? What is the quantile (25%, 50%, 75%) for these numbers? Which gene has the greatest number of transcript?

```
d %>% group_by(gene_id) %>% count(type) %>% filter(type=="transcript")
```

```
## # A tibble: 60,603 x 3
## # Groups:   gene_id [60,603]
##   gene_id          type      n
##   <chr>          <fct>    <int>
## 1 ENSG000000000003.14 transcript     3
## 2 ENSG000000000005.6  transcript     1
## 3 ENSG000000000419.12 transcript     2
```

```
## 4 ENSG00000000457.14 transcript 3
## 5 ENSG00000000460.17 transcript 5
## 6 ENSG00000000938.13 transcript 3
## 7 ENSG00000000971.15 transcript 3
## 8 ENSG00000001036.13 transcript 1
## 9 ENSG00000001084.13 transcript 5
## 10 ENSG00000001167.14 transcript 2
## # ... with 60,593 more rows
```

```
d %>% group_by(gene_id) %>% count(type) %>% filter(type=="transcript") %>% pull() %>% mean(.)
```

```
## [1] 1.7861
```

```
d %>% group_by(gene_id) %>% count(type) %>% filter(type=="transcript") %>% pull() %>%
  quantile(.,c(0.25,0.50,0.75))
```

```
## 25% 50% 75%
##    1    1    2
```

```
d %>% group_by(gene_id) %>% count(type) %>% filter(type=="transcript") %>% ungroup() %>% top_n(1,n)
```

```
## # A tibble: 1 x 3
##   gene_id      type      n
##   <chr>      <fct>   <int>
## 1 ENSG00000109339.22 transcript 87
```

2. Compute the number of transcripts per gene among gene biotypes. For example, compare the number of transcript per gene between protein-coding genes, long noncoding genes, pseudogenes.

```
d %>% group_by(gene_type) %>% count(type) %>% filter(type=="transcript")
```

```
## # A tibble: 40 x 3
## # Groups:   gene_type [40]
##   gene_type      type      n
##   <chr>      <fct>   <int>
## 1 IG_C_gene      transcript 14
## 2 IG_C_pseudogene transcript  9
## 3 IG_D_gene      transcript 37
## 4 IG_J_gene      transcript 18
## 5 IG_J_pseudogene transcript  3
## 6 IG_pseudogene   transcript  1
## 7 IG_V_gene      transcript 144
## 8 IG_V_pseudogene transcript 188
## 9 lncRNA          transcript 24993
## 10 miRNA          transcript 1881
## # ... with 30 more rows
```

3. Final task is to compute the number of transcripts per gene per chromosome.

```
d %>% group_by(seqnames) %>% count(type) %>% filter(type=="transcript")
```

```
## # A tibble: 25 x 3
## # Groups:   seqnames [25]
##   seqnames type      n
##   <fct>    <fct>    <int>
## 1 chr1     transcript  9827
## 2 chr2     transcript  7432
## 3 chr3     transcript  6157
## 4 chr4     transcript  4662
## 5 chr5     transcript  5203
## 6 chr6     transcript  5455
## 7 chr7     transcript  5292
## 8 chr8     transcript  4350
## 9 chr9     transcript  3949
## 10 chr10    transcript  4157
## # ... with 15 more rows
```

### 3.2. Gene length in the GENCODE

1. What is the average length of human genes?

```
d %>% select(width) %>% pull() %>% mean(.)
```

```
## [1] 4069.518
```

2. Is the distribution of gene length differed by autosomal and sex chromosomes? Please calculate the quantiles (0%, 25%, 50%, 75%, 100%) of the gene length for each group.

```
d %>% filter(seqnames!="chrM") %>%
  mutate(chr = ifelse(seqnames %in% c("chrX","chrY"),"sex_chr","autosomal_chr")) %>%
  group_by(chr) %>%
  summarize(quantiles=quantile(width,c(0,0.25,0.50,0.75,1)))
```

## `summarise()` has grouped output by 'chr'. You can override using the `.groups` argument.

```
## # A tibble: 10 x 2
## # Groups:   chr [2]
##   chr      quantiles
##   <chr>         <dbl>
## 1 autosomal_chr      1
## 2 autosomal_chr     80
## 3 autosomal_chr    129
## 4 autosomal_chr    222
## 5 autosomal_chr 2473537
## 6 sex_chr            1
## 7 sex_chr           78
## 8 sex_chr          127
## 9 sex_chr          230
## 10 sex_chr       2241765
```

3. Is the distribution of gene length differed by gene biotype? Please calculate the quantiles (0%, 25%, 50%, 75%, 100%) of the gene length for each group.

```
d %>% group_by(gene_type) %>% summarize(quantiles=quantile(width,c(0,0.25,0.50,0.75,1)))
```

## `summarise()` has grouped output by 'gene\_type'. You can override using the `.groups` argument.

```
## # A tibble: 200 x 2
## # Groups:   gene_type [40]
##   gene_type      quantiles
##   <chr>         <dbl>
## 1 IG_C_gene         3
## 2 IG_C_gene        92
## 3 IG_C_gene       312.
## 4 IG_C_gene       336
## 5 IG_C_gene      8914
## 6 IG_C_pseudogene   34
## 7 IG_C_pseudogene  293
## 8 IG_C_pseudogene  316
## 9 IG_C_pseudogene  424
## 10 IG_C_pseudogene 5211
## # ... with 190 more rows
```

### 3.3. Transcript support levels (TSL)

The GENCODE TSL provides a consistent method of evaluating the level of support that a GENCODE transcript annotation is actually expressed in humans.

1. With transcript, how many transcripts are categorized for each TSL?

```
d %>% group_by(transcript_support_level) %>% count(type) %>% filter(type=="transcript")
```

```
## # A tibble: 7 x 3
## # Groups:   transcript_support_level [7]
##   transcript_support_level type      n
##   <chr>                  <fct>   <int>
## 1 1                      transcript 31801
## 2 2                      transcript 13372
## 3 3                      transcript  7228
## 4 4                      transcript  2245
## 5 5                      transcript 13674
## 6 NA                    transcript 27843
## 7 <NA>                  transcript 12080
```

2. From the first question, please count the number of transcript for each TSL by gene biotype.

```
d %>% group_by(transcript_support_level, gene_type) %>% count(type) %>% filter(type=="transcript")
```

```
## # A tibble: 91 x 4
## # Groups:   transcript_support_level, gene_type [91]
```

```
## transcript_support_level gene_type type n
## <chr> <chr> <fct> <int>
## 1 1 IG_C_gene transcript 1
## 2 1 lncRNA transcript 1620
## 3 1 polymorphic_pseudogene transcript 30
## 4 1 protein_coding transcript 29783
## 5 1 transcribed_processed_pseudogene transcript 42
## 6 1 transcribed_unitary_pseudogene transcript 58
## 7 1 transcribed_unprocessed_pseudogene transcript 267
## 8 2 lncRNA transcript 2970
## 9 2 polymorphic_pseudogene transcript 3
## 10 2 protein_coding transcript 10104
## # ... with 81 more rows
```

3. From the first question, please count the number of transcript for each TSL by source.

```
d %>% group_by(transcript_support_level,source) %>% count(type) %>% filter(type=="transcript")
```

```
## # A tibble: 14 x 4
## # Groups: transcript_support_level, source [14]
## transcript_support_level source type n
## <chr> <fct> <fct> <int>
## 1 1 HAVANA transcript 29434
## 2 1 ENSEMBL transcript 2367
## 3 2 HAVANA transcript 12052
## 4 2 ENSEMBL transcript 1320
## 5 3 HAVANA transcript 6964
## 6 3 ENSEMBL transcript 264
## 7 4 HAVANA transcript 2116
## 8 4 ENSEMBL transcript 129
## 9 5 HAVANA transcript 10157
## 10 5 ENSEMBL transcript 3517
## 11 NA HAVANA transcript 19962
## 12 NA ENSEMBL transcript 7881
## 13 <NA> HAVANA transcript 11901
## 14 <NA> ENSEMBL transcript 179
```

### 3.4. CCDS in the GENCODE

1. With gene, please create a data frame with the columns - gene\_id, gene\_name, hgnc\_id, gene\_type, chromosome, start, end, and strand. Then, please create new columns for presence of hgnc and ccds. For example, you can put 1 in the column isHgnc, if hgnc annotation is available, or 0 if not. Then, you can put 1 in the column isCCDS, if ccds annotation is available, or 0 if not.

```
d %>% mutate(isHgnc=ifelse(is.na(hgnc_id),0,1),isCCDS=ifelse(is.na(ccdsid),0,1)) %>%
  select(gene_id,gene_name,hgnc_id,gene_type,seqnames,start,end,strand,isHgnc,isCCDS) %>%
  head()
```

```
## gene_id gene_name hgnc_id gene_type
## 1 ENSG00000223972.5 DDX11L1 HGNC:37102 transcribed_unprocessed_pseudogene
## 2 ENSG00000223972.5 DDX11L1 HGNC:37102 transcribed_unprocessed_pseudogene
## 3 ENSG00000223972.5 DDX11L1 HGNC:37102 transcribed_unprocessed_pseudogene
```

```
## 4 ENSG00000223972.5 DDX11L1 HGNC:37102 transcribed_unprocessed_pseudogene
## 5 ENSG00000223972.5 DDX11L1 HGNC:37102 transcribed_unprocessed_pseudogene
## 6 ENSG00000223972.5 DDX11L1 HGNC:37102 transcribed_unprocessed_pseudogene
##   seqnames start   end strand isHgnc isCCDS
## 1   chr1 11869 14409      +      1      0
## 2   chr1 11869 14409      +      1      0
## 3   chr1 11869 12227      +      1      0
## 4   chr1 12613 12721      +      1      0
## 5   chr1 13221 14409      +      1      0
## 6   chr1 12010 13670      +      1      0
```

2. Please count the number of hgnc by gene biotypes.

```
d %>% group_by(gene_type) %>% count(hgnc_id)
```

```
## # A tibble: 37,565 x 3
## # Groups:   gene_type [40]
##   gene_type hgnc_id      n
##   <chr>     <chr>    <int>
## 1 IG_C_gene HGNC:5478    14
## 2 IG_C_gene HGNC:5479    14
## 3 IG_C_gene HGNC:5480    16
## 4 IG_C_gene HGNC:5522    16
## 5 IG_C_gene HGNC:5525    16
## 6 IG_C_gene HGNC:5526    16
## 7 IG_C_gene HGNC:5527    22
## 8 IG_C_gene HGNC:5528    16
## 9 IG_C_gene HGNC:5541    16
## 10 IG_C_gene HGNC:5716     6
## # ... with 37,555 more rows
```

3. Please count the number of hgnc by level. Please note that level in this question is not TSL. Please find information in this link: 1 (verified loci), 2 (manually annotated loci), 3 (automatically annotated loci).

```
d %>% group_by(level) %>% count(hgnc_id)
```

```
## # A tibble: 49,081 x 3
## # Groups:   level [3]
##   level hgnc_id      n
##   <chr> <chr>    <int>
## 1 1      HGNC:100     26
## 2 1      HGNC:10001    16
## 3 1      HGNC:10002    37
## 4 1      HGNC:10007     1
## 5 1      HGNC:10008    30
## 6 1      HGNC:1001     22
## 7 1      HGNC:10010     3
## 8 1      HGNC:1004     18
## 9 1      HGNC:10049     3
## 10 1      HGNC:10055     1
## # ... with 49,071 more rows
```

### 3.5. Transcripts in the GENCODE

1. Which gene has the largest number of transcripts?

```
d %>% group_by(gene_id) %>% count(type) %>% filter(type=="transcript") %>% ungroup() %>% top_n(1,n)
```

```
## # A tibble: 1 x 3
##   gene_id      type      n
##   <chr>      <fct>    <int>
## 1 ENSG00000109339.22 transcript    87
```

2. Please calculate the quantiles (0%, 25%, 50%, 75%, 100%) of the gene length for protein coding genes and long noncoding genes.

```
d %>% filter(gene_type %in% c("protein_coding","lncRNA")) %>% group_by(gene_type) %>%
  summarize(quantiles=quantile(width,c(0,0.25,0.50,0.75,1)))
```

```
## `summarise()` has grouped output by 'gene_type'. You can override using the `.groups` argument.
```

```
## # A tibble: 10 x 2
## # Groups:   gene_type [2]
##   gene_type      quantiles
##   <chr>          <dbl>
## 1 lncRNA          1
## 2 lncRNA        126
## 3 lncRNA        339
## 4 lncRNA       2658
## 5 lncRNA     1375317
## 6 protein_coding    1
## 7 protein_coding   76
## 8 protein_coding  123
## 9 protein_coding  193
## 10 protein_coding 2473537
```

3. Please count the number of transcripts by chromosomes.

```
d %>% group_by(seqnames) %>% count(type) %>% filter(type=="transcript")
```

```
## # A tibble: 25 x 3
## # Groups:   seqnames [25]
##   seqnames type      n
##   <fct>    <fct>    <int>
## 1 chr1    transcript  9827
## 2 chr2    transcript  7432
## 3 chr3    transcript  6157
## 4 chr4    transcript  4662
## 5 chr5    transcript  5203
## 6 chr6    transcript  5455
## 7 chr7    transcript  5292
## 8 chr8    transcript  4350
## 9 chr9    transcript  3949
## 10 chr10   transcript  4157
## # ... with 15 more rows
```

### 3.6. Autosomal vs. Sex chromosomes.

1. Please calculate the number of genes per chromosome.

```
d %>% group_by(seqnames) %>% count()
```

```
## # A tibble: 25 x 2
## # Groups:   seqnames [25]
##   seqnames      n
##   <fct>      <int>
## 1 chr1      167739
## 2 chr2      130271
## 3 chr3      107787
## 4 chr4       72009
## 5 chr5       77967
## 6 chr6       86626
## 7 chr7       85829
## 8 chr8       61597
## 9 chr9       66590
## 10 chr10      72048
## # ... with 15 more rows
```

2. Please compare the number of genes between autosomal and sex chromosome (Mean, Median).

```
d3.5.2 <- d %>% filter(seqnames!="chrM") %>%
  mutate(chr = ifelse(seqnames %in% c("chrX","chrY"), "sex_chr", "autosomal_chr"))
d3.5.2 %>% filter(chr=="autosomal_chr") %>% group_by(seqnames) %>% count() %>% pull() %>% mean()
```

```
## [1] 76542.14
```

```
d3.5.2 %>% filter(chr=="autosomal_chr") %>% group_by(seqnames) %>% count() %>% pull() %>% median()
```

```
## [1] 73309.5
```

```
d3.5.2 %>% filter(chr=="sex_chr") %>% group_by(seqnames) %>% count() %>% pull() %>% mean()
```

```
## [1] 36216
```

```
d3.5.2 %>% filter(chr=="sex_chr") %>% group_by(seqnames) %>% count() %>% pull() %>% median()
```

```
## [1] 36216
```

3. Please divide the genes into groups 'protein coding' and 'long noncoding', and then compare the number of genes in each chromosomes within groups.

```
d %>% filter(gene_type %in% c("protein_coding", "lncRNA")) %>% group_by(gene_type) %>% count(seqnames)
```



```
## # A tibble: 49 x 3
## # Groups:   gene_type [2]
##   gene_type seqnames      n
##   <chr>      <fct>    <int>
## 1 lncRNA      chr1     10317
## 2 lncRNA      chr2     10361
## 3 lncRNA      chr3       7647
## 4 lncRNA      chr4       6546
## 5 lncRNA      chr5       7415
## 6 lncRNA      chr6       6505
## 7 lncRNA      chr7       5670
## 8 lncRNA      chr8       6680
## 9 lncRNA      chr9       4013
## 10 lncRNA     chr10      5209
## # ... with 39 more rows
```